



MONASH University

Information Technology

FIT2100 Operating Systems

REAL-TIME PROCESS SCHEDULING

Week 6 - Supplementary

Semester 2 2024

(Reading: Tanenbaum, Chapter 2)

Dr Charith Jayasekara

Faculty of Information Technology

© 2024 Monash University

SCHEDULING IN REAL-TIME SYSTEM



SCHEDULING IN REAL-TIME SYSTEM





- Timeliness is of critical importance
- Categories
 - **Hard real time**
 - ❑ there are absolute deadlines that must be met
 - **Soft real time**
 - ❑ meaning that missing an occasional deadline is undesirable, but tolerable
 - In both categories, real-time behavior is achieved by dividing the program into a number of processes, the execution time of which is predictable and known in advance. These processes are generally short lived and can run to completion in well under a second. When an external event is detected, it is the job of the scheduler to schedule the processes in such a way that all deadlines are met.

SCHEDULING IN REAL-TIME SYSTEM

- Periodic or aperiodic
 - Periodic - meaning they occur at regular intervals
 - Aperiodic - meaning they occur unpredictably (e.g. randomly or chaotically)


SCHEDULING IN REAL-TIME SYSTEMS

100% PERIODIC



100% APERIODIC
Periodic or aperiodic

SCHEDULING IN REAL-TIME SYSTEMS



periodic or aperiodic

SCHEDULING IN REAL-TIME SYSTEMS

- Schedulable satisfies

- A system may have to respond to multiple periodic event streams. Depending on how much time each event requires for processing, handling all of them may not even be possible.
- For example, if there are m periodic events and event i occurs with period P_i and requires C_i sec of CPU time to handle each event, then the load can be handled only if A real-time system that meets this criterion is said to be schedulable.
- This means it can actually be implemented.
- A process that fails to meet this test cannot be scheduled because the total amount of CPU time the processes want collectively is more than the CPU can deliver

$$\sum_{i=1}^m \frac{C_i}{P_i} \leq 1$$

SCHEDULING IN REAL-TIME SYSTEMS

- Schedulable satisfies
 - Example, consider a soft real-time system with three periodic events, with periods of 100, 200, and 500 msec, respectively. If these events require 50, 30, and 100 msec of CPU time per event, respectively.
 - The system is schedulable because $0.5 + 0.15 + 0.2 < 1$.
 - If a fourth event with a period of 1 sec is added, the system will remain schedulable as long as this event does not need more than 150 msec of CPU time per event.
 - Implicit in this calculation is the assumption that the context-switching overhead is so small that it can be ignored.
 - If context-switching time is non-negligible, it must be included in the calculation.

$$\sum_{i=1}^m \frac{C_i}{P_i} \leq 1$$

SCHEDULING IN REAL-TIME SYSTEMS

STATIC VS DYNAMIC SCHEDULING

- **Static**

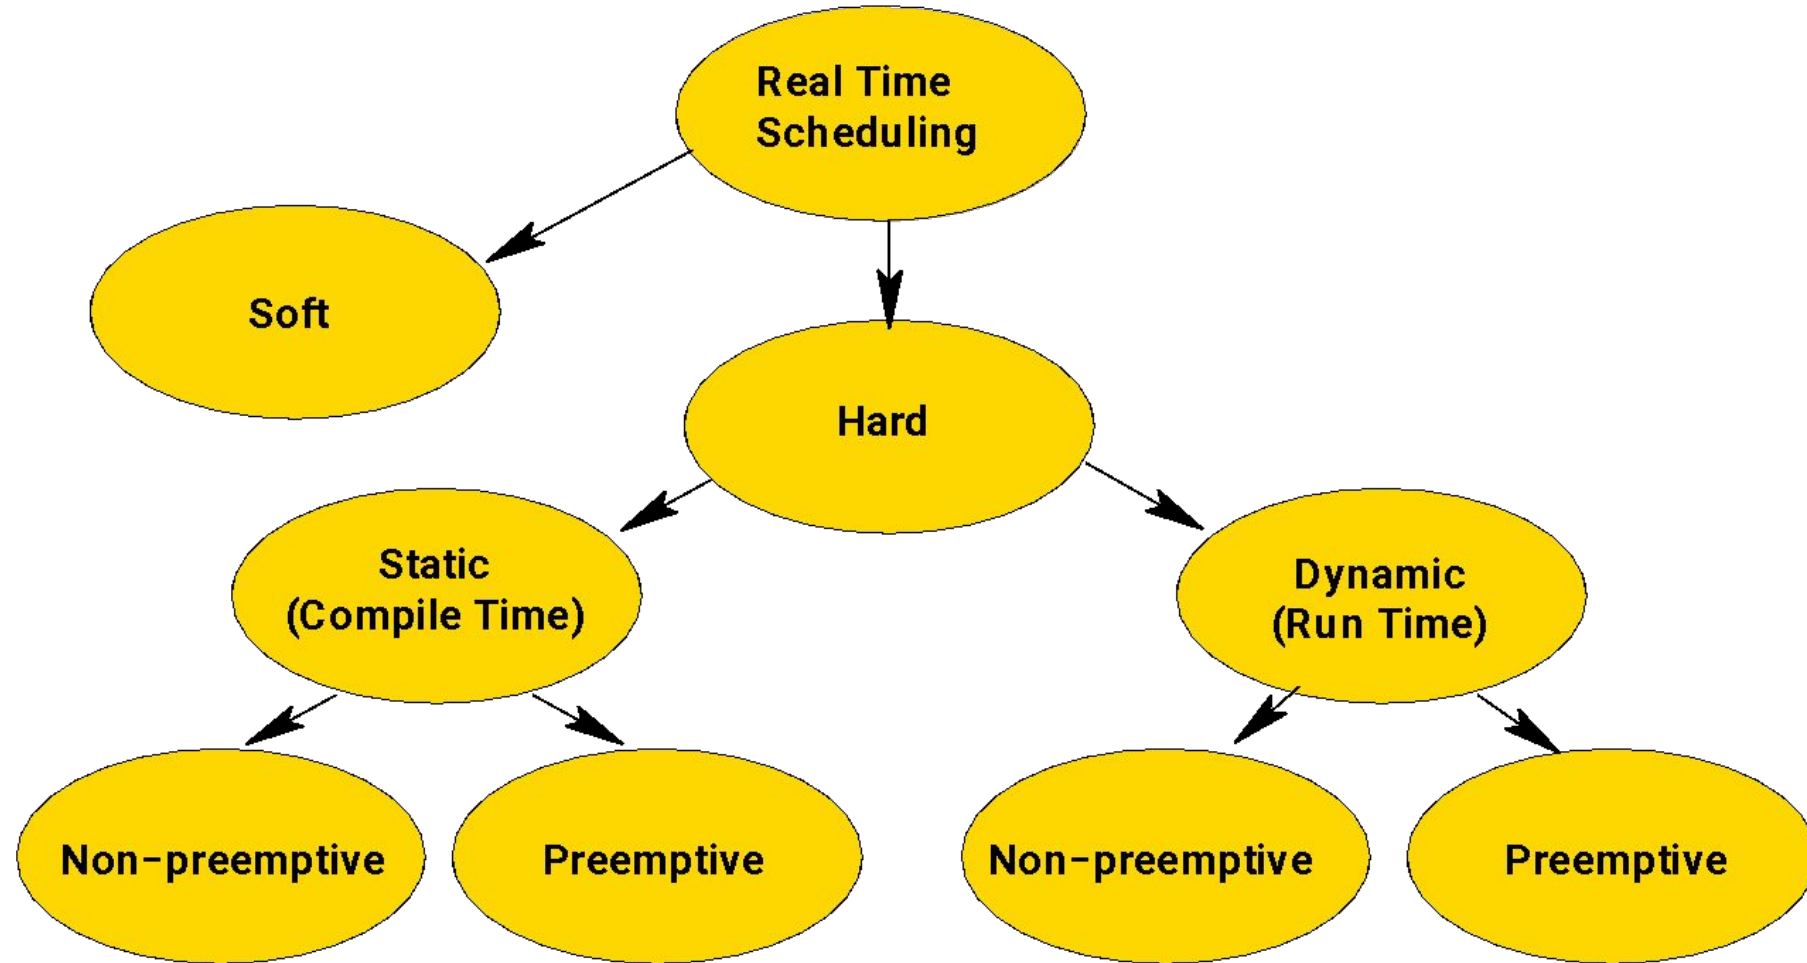
- Make their scheduling decisions before the system starts running.
- Static scheduling works only when there is perfect information available in advance about the work to be done and the deadlines that have to be met.

- **Dynamic**

- Make their scheduling decisions at run time, after execution has started.
- Dynamic scheduling algorithms do not have the restrictions that apply to static scheduling.

SCHEDULING IN REAL-TIME SYSTEMS

REAL TIME SCHEDULING TAXONOMY



REAL-TIME OPERATING SYSTEMS

Examples of Commonly Used RTOS Types

- **VxWorks** (Wind River Systems), supports x86, MIPS, ARM, PowerPC, SuperH, RISC-V
- **Nucleus** (Siemens EDA), supports ARM, MIPS, PowerPC, RISC-V (replaced VRTX)
- **LynxOS** (Lynx Software Technologies), supports x86, 68k, ARM, PowerPC
- **INTEGRITY** (Green Hills Software), supports ~40 architectures
- **QNX Neutrino** (Blackberry), supports x86-64, ARM32, ARM64
- **Windows IoT** (Microsoft), family of embedded OS variants of Windows
- **FreeRTOS** (Freertos.org), supports 40+ microcontroller architectures
- **RTAI** (RTAI.org), Linux kernel extension, supports x86, PowerPC, ARM, MIPS
- NB in the embedded and IoT sectors, ARM and i86 are common but not exclusive!

ASSESSING RTOS PERFORMANCE



ASSESSING RTOS PERFORMANCE

METRICS

- **Turnaround Time**

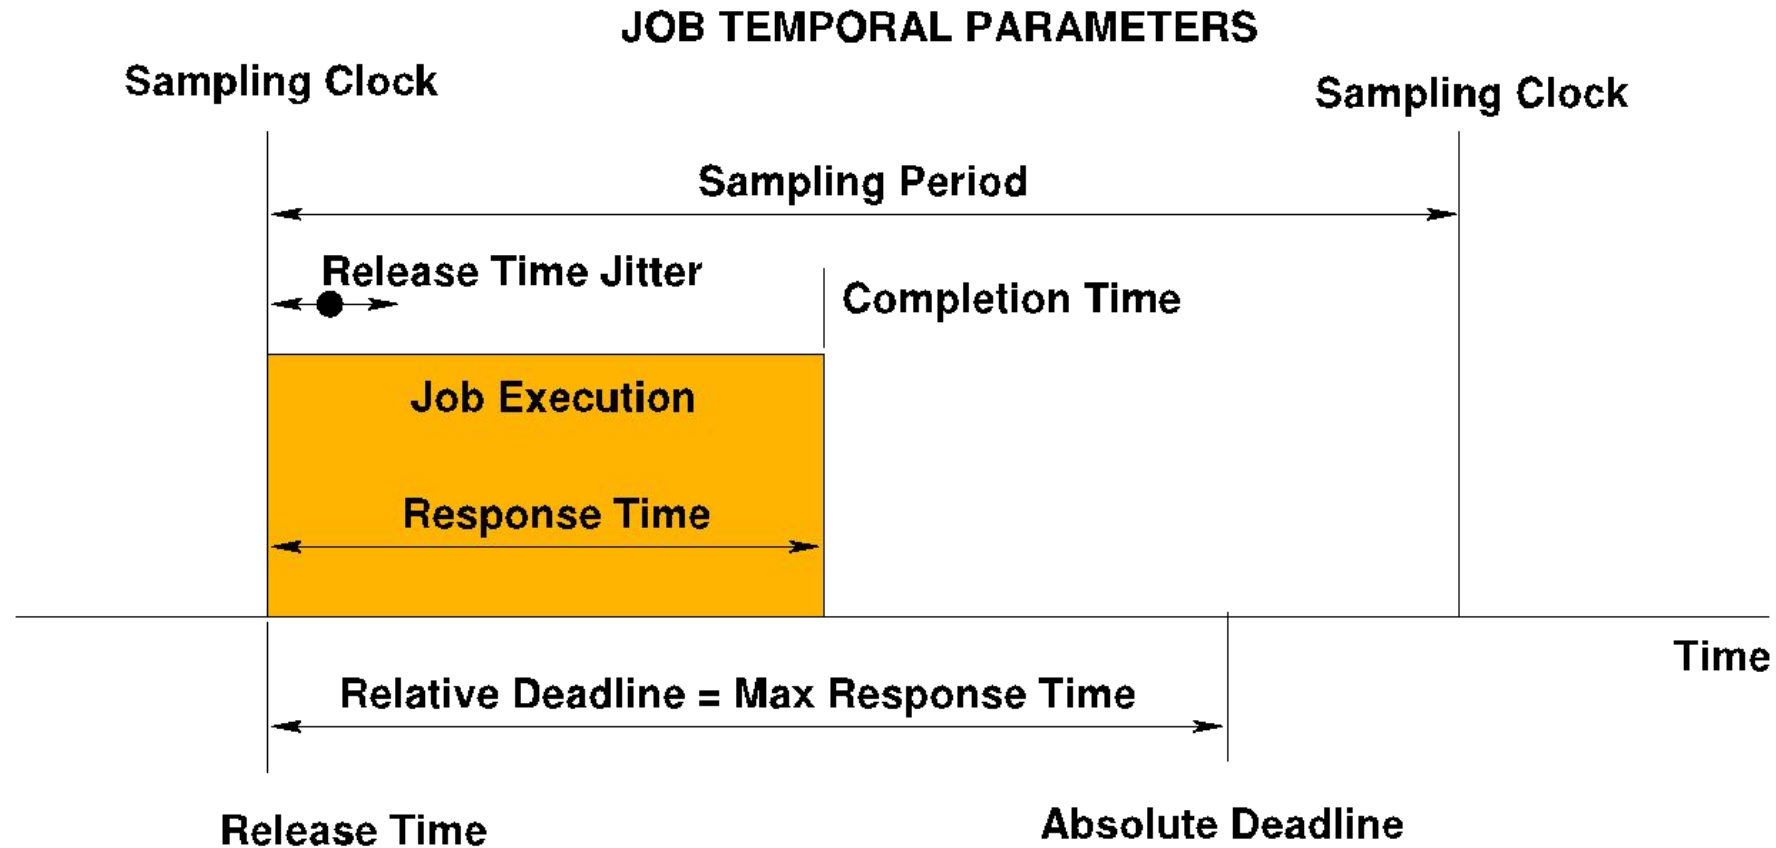
- Time interval from entry to completion
- The total time a process spends either running or waiting, until it completes.

□ $\text{Total time} = \text{Completion time} - \text{Arrival time}$

- **Response Time (Interrupt Latency)**

- Average time between arrival time, and the process beginning to respond (i.e. returning to the RUNNING state)
- Critically important for many hard real-time processes

ASSESSING PERFORMANCE



If (Response Time) > (Max Response Time) => Hard Real Time Job Fails!
Example: A Periodic Hard Real-Time Process (e.g. flight control system)

ASSESSING RTOS PERFORMANCE

METRICS

- Throughput
 - The average rate at which processes are passing through the system from entry to completion
 - Can be measured as number of processes per time taken, or its inverse.
 - E.g. a system admits and completes 12 processes in 6 seconds.
 - $\text{Throughput} = 12 \text{ processes} / 6 \text{ sec} = 2 \text{ proc/sec}$
 - Or inverse: $6 \text{ sec} / 12 \text{ proc} = 0.5 \text{ seconds per process}$

REAL-TIME SCHEDULING ALGORITHMS



SCHEDULING IN REAL-TIME SYSTEMS

Earliest-Deadline-First (EDF) Algorithm

- A way to assign priorities to jobs is on the basis of their deadlines.
- Earliest-Deadline-First (EDF) algorithm is based on the priority assignment whereby the earlier the deadline, the higher the priority.
- This algorithm is important because it is optimal when used to schedule jobs on a processor when preemption is allowed and there is no contention for resources
- When preemption is allowed and jobs do not contend for resources, the EDF algorithm can produce a feasible schedule of a set of jobs J with arbitrary release times and deadlines on a processor, if and only if J has feasible schedules. resources.

SCHEDULING IN REAL-TIME SYSTEMS

EXAMPLE: Earliest-Deadline-First (EDF) Algorithm (Zagan and Gaitan, 2016)

TABLE II

THE PARAMETERS OF A SET OF FIVE TASKS

	a_i	C_i	D_i
τ_5	12	5	18
τ_4	7	8	24
τ_3	4	4	10
τ_2	1	3	11
τ_1	0	2	4

τ_i – task identifier

a_i – task arrival time

C_i – task compute duration

D_i – task deadline

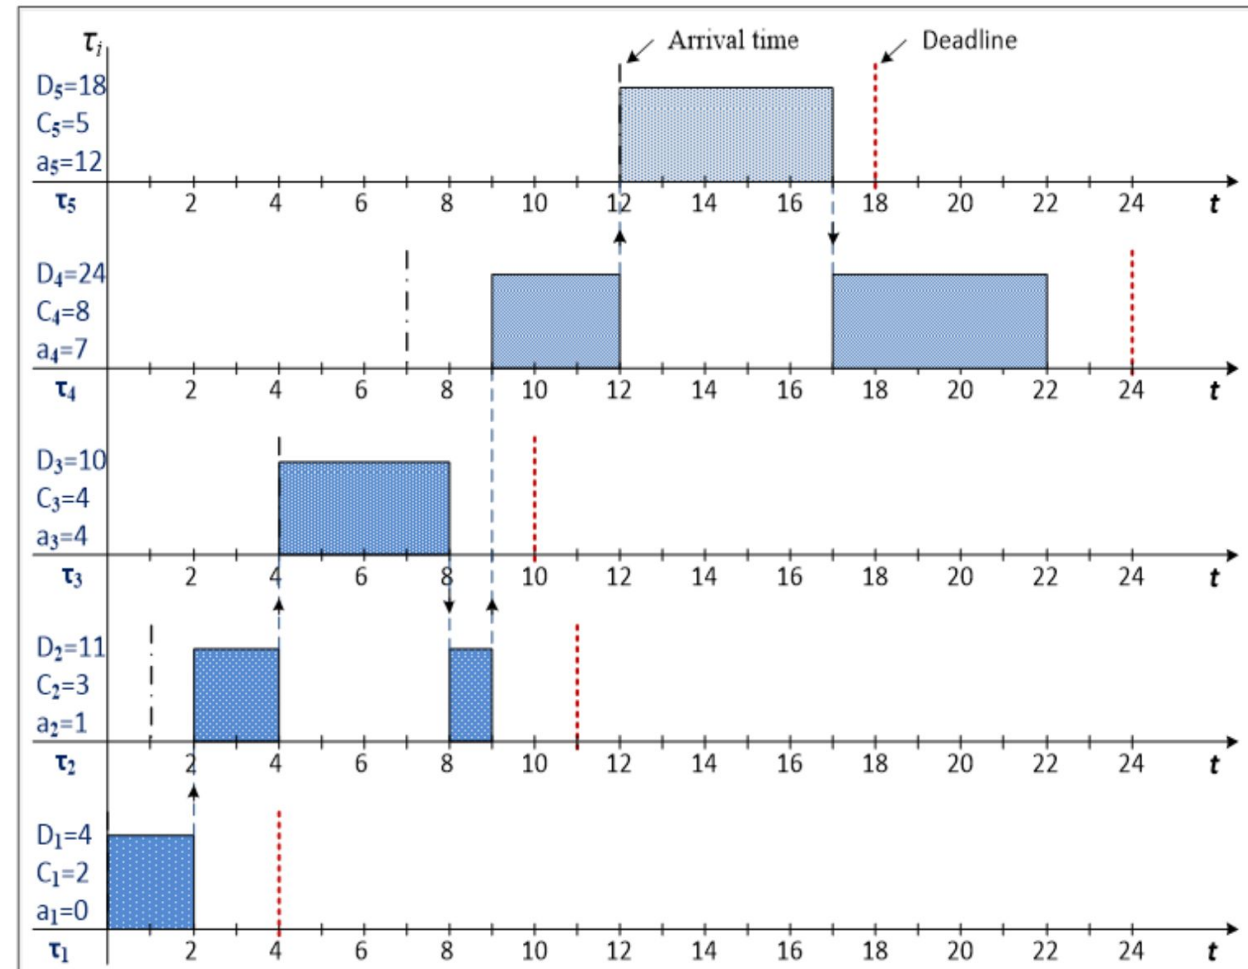


Fig. 1. A scheduling example using the EDF algorithm.

SCHEDULING IN REAL-TIME SYSTEMS

EXAMPLE: Earliest-Deadline-First (EDF) Algorithm (Zagan and Gaitan, 2016)

- Cited: “Fig. 1 shows an example of scheduling a set of five tasks using the EDF algorithm. At moment $t = 0$, task τ_1 enters execution, and at moment $t = 1$, task τ_2 cannot interrupt τ_1 because $D_1 < D_2$. Task τ_1 completes execution at time moment $t = 2$, and at moment $t = 4$, when τ_2 is being executed, task τ_3 interrupts τ_2 because $D_3 < D_2$. To be noted that at time moment $t = 7$, task τ_4 does not interrupt τ_3 because $D_3 < D_4$. When τ_3 completes execution, the CPU is assigned to task τ_2 . Task τ_4 is executed at moment $t = 9$, but it is interrupted at $t = 12$ by τ_5 , because the last one has a lower deadline. Task τ_4 re-enters in execution at moment $t = 17$, when τ_5 completes its own.”

SCHEDULING IN REAL-TIME SYSTEMS

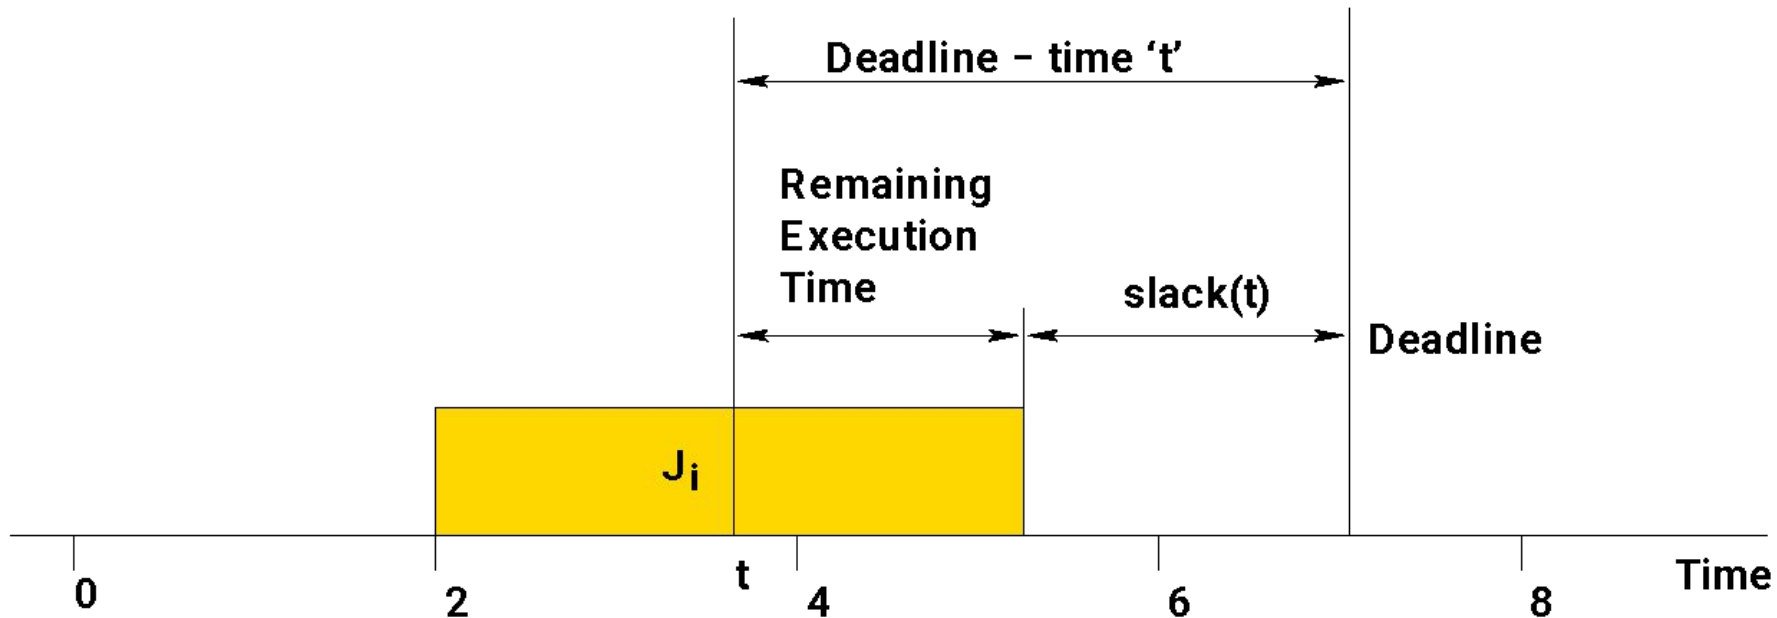
Least-Slack-Time-First (LST) Algorithm

- Another algorithm optimal for scheduling preemptive jobs on one processor is the Least-Slack-Time-First (LST), also called Minimum-Laxity-First (MLF) algorithm.
- At any time t , the slack (or laxity) of a job with deadline d , is equal to $d - t$ minus the time required to complete the remainder of the job.
- $\text{slack}(t) = (\text{deadline} - t) - (\text{execution time remaining}(t))$
- $\text{slack}(t) = \text{deadline} - t - \text{etr}(t)$

REAL-TIME SCHEDULING ALGORITHMS

LEAST SLACK TIME FIRST ALGORITHM

$\text{slack}(t) = d_i - t - \text{etr}_i(t)$; etr = execution time remaining
slack is re-calculated when a job is released or a job completes



REAL-TIME SCHEDULING ALGORITHMS

LEAST SLACK TIME FIRST ALGORITHM

Execution Time (Feasible Interval)



$\text{slack}(t) = d_i - t - \text{etr}_i(t)$; etr = execution time remaining
slack is re-calculated when a job is released or a job completes

$J_1, t = 0: r_1 = 0, d_1 = 6, \text{etr}_1 = 3 \Rightarrow \text{slack time} = 6 - 0 - 3 = 3 @ t = 0$


$J_1, t = 2: r_1 = 0, d_1 = 6, \text{etr}_1 = 1 \Rightarrow \text{slack time} = 6 - 2 - 1 = 3 @ t = 2$

$J_3, t = 2: r_3 = 2, d_3 = 7, \text{etr}_3 = 2 \Rightarrow \text{slack time} = 7 - 2 - 2 = 3 @ t = 2$

$J_1, t = 4: r_1 = 0, d_1 = 6, \text{etr}_1 = 1 \Rightarrow \text{slack time} = 6 - 4 - 1 = 1 @ t = 4$

NB J_2 cannot be calculated since it is released at $t = 5$

Pre-emption (arbitrary since identical slack magnitude)



REAL-TIME SCHEDULING ALGORITHMS

LEAST SLACK TIME FIRST ALGORITHM

Execution Time (Feasible Interval)

J_1 3 (0,6]

J_2 2 (5,8]

J_3 2 (2,7]




$t = 0$: $r_1 = 0$, $d_1 = 6$, $e_1 = 3 \Rightarrow$ slack time = 3 @ $t = 0$

i.e. slack time for $J_1 = 6 - t - (3 - t)$ while J_1 is executing

$t = 4$: $r_1 = 0$, $d_1 = 6$, remaining $e_1 = 1 \Rightarrow$ slack time = 1 @ $t = 4$

$t = 4$: $r_2 = 5$, $d_2 = 8$, $e_2 = 2 \Rightarrow$ slack time = $4 + 2 = 6$ @ $t = 4$

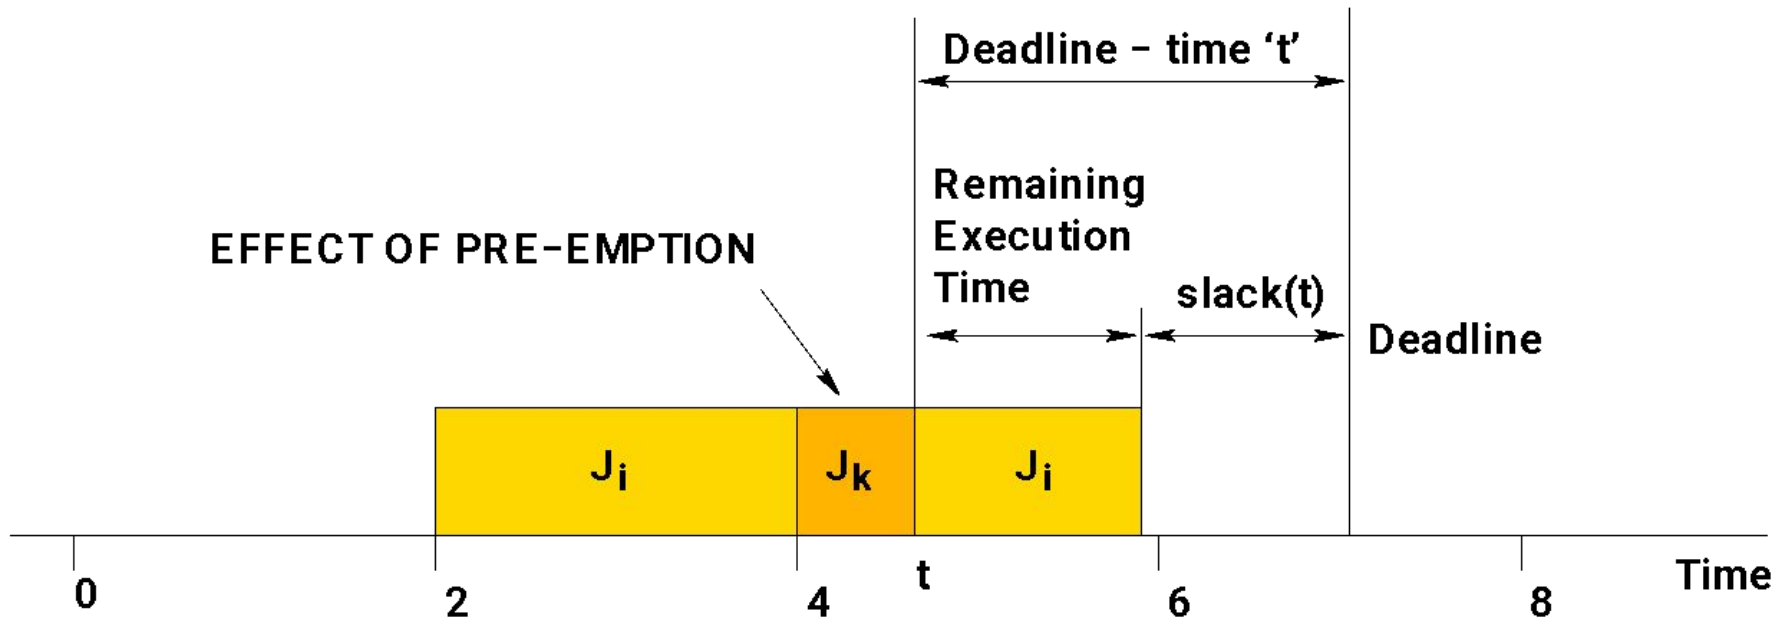
Pre-emption



REAL-TIME SCHEDULING ALGORITHMS

LEAST SLACK TIME FIRST ALGORITHM

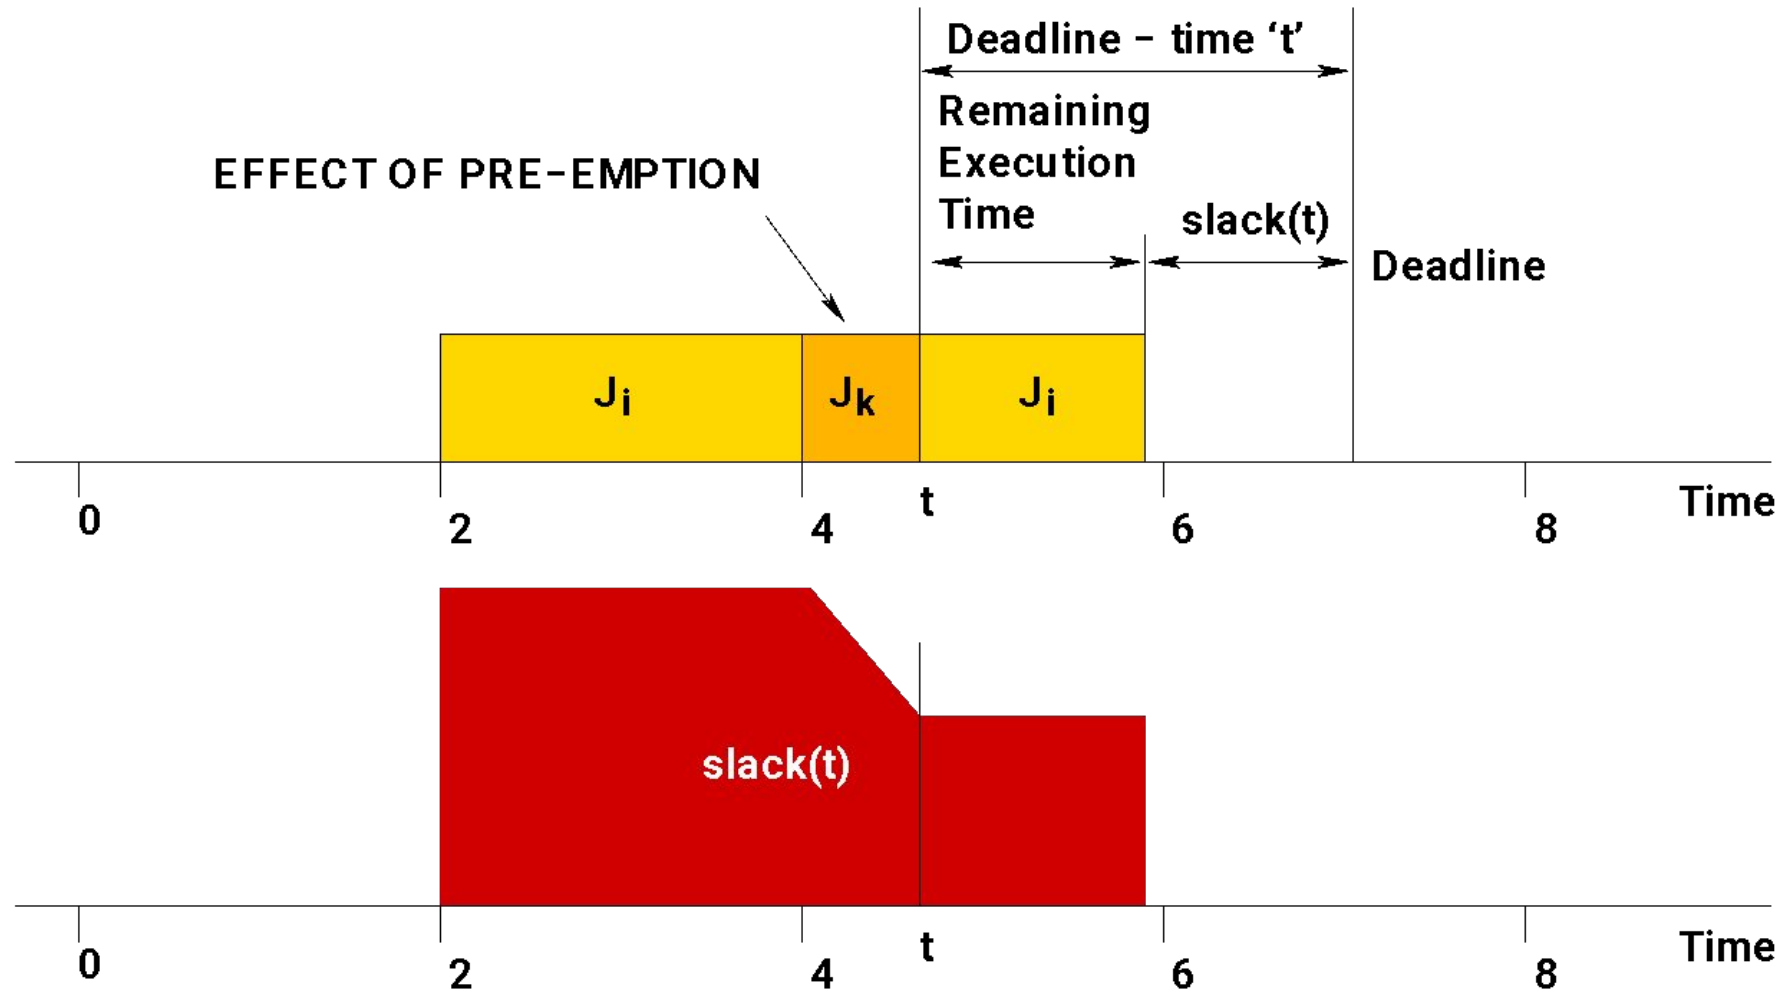
$\text{slack}(t) = d_i - t - \text{etr}_i(t)$; etr = execution time remaining
slack is re-calculated when a job is released or a job completes



Every time a job is pre-empted its slack time decreases

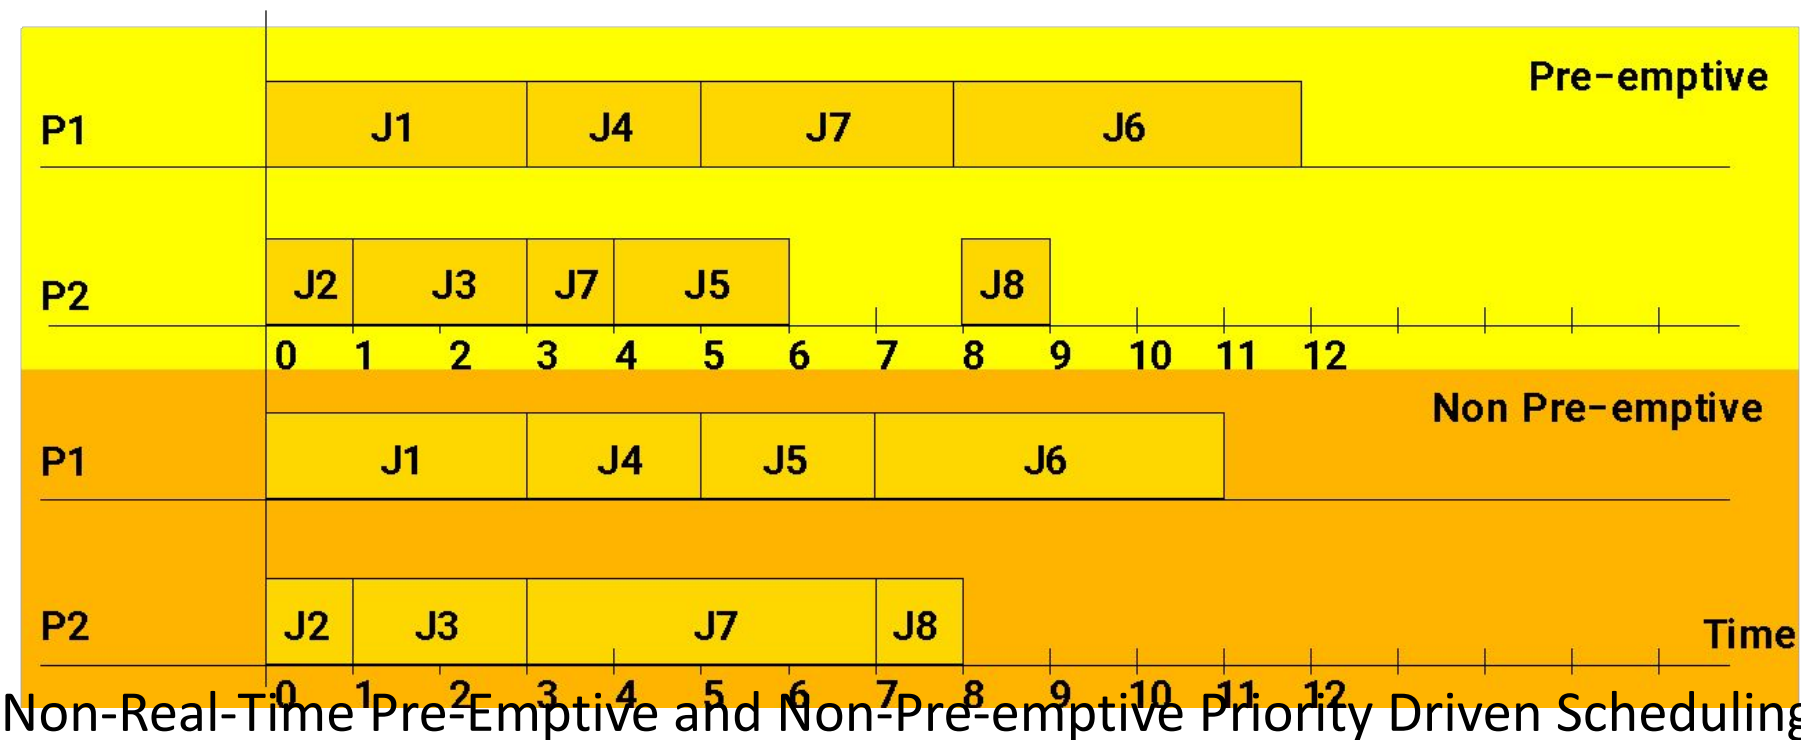

REAL-TIME SCHEDULING ALGORITHMS

LEAST SLACK TIME FIRST ALGORITHM

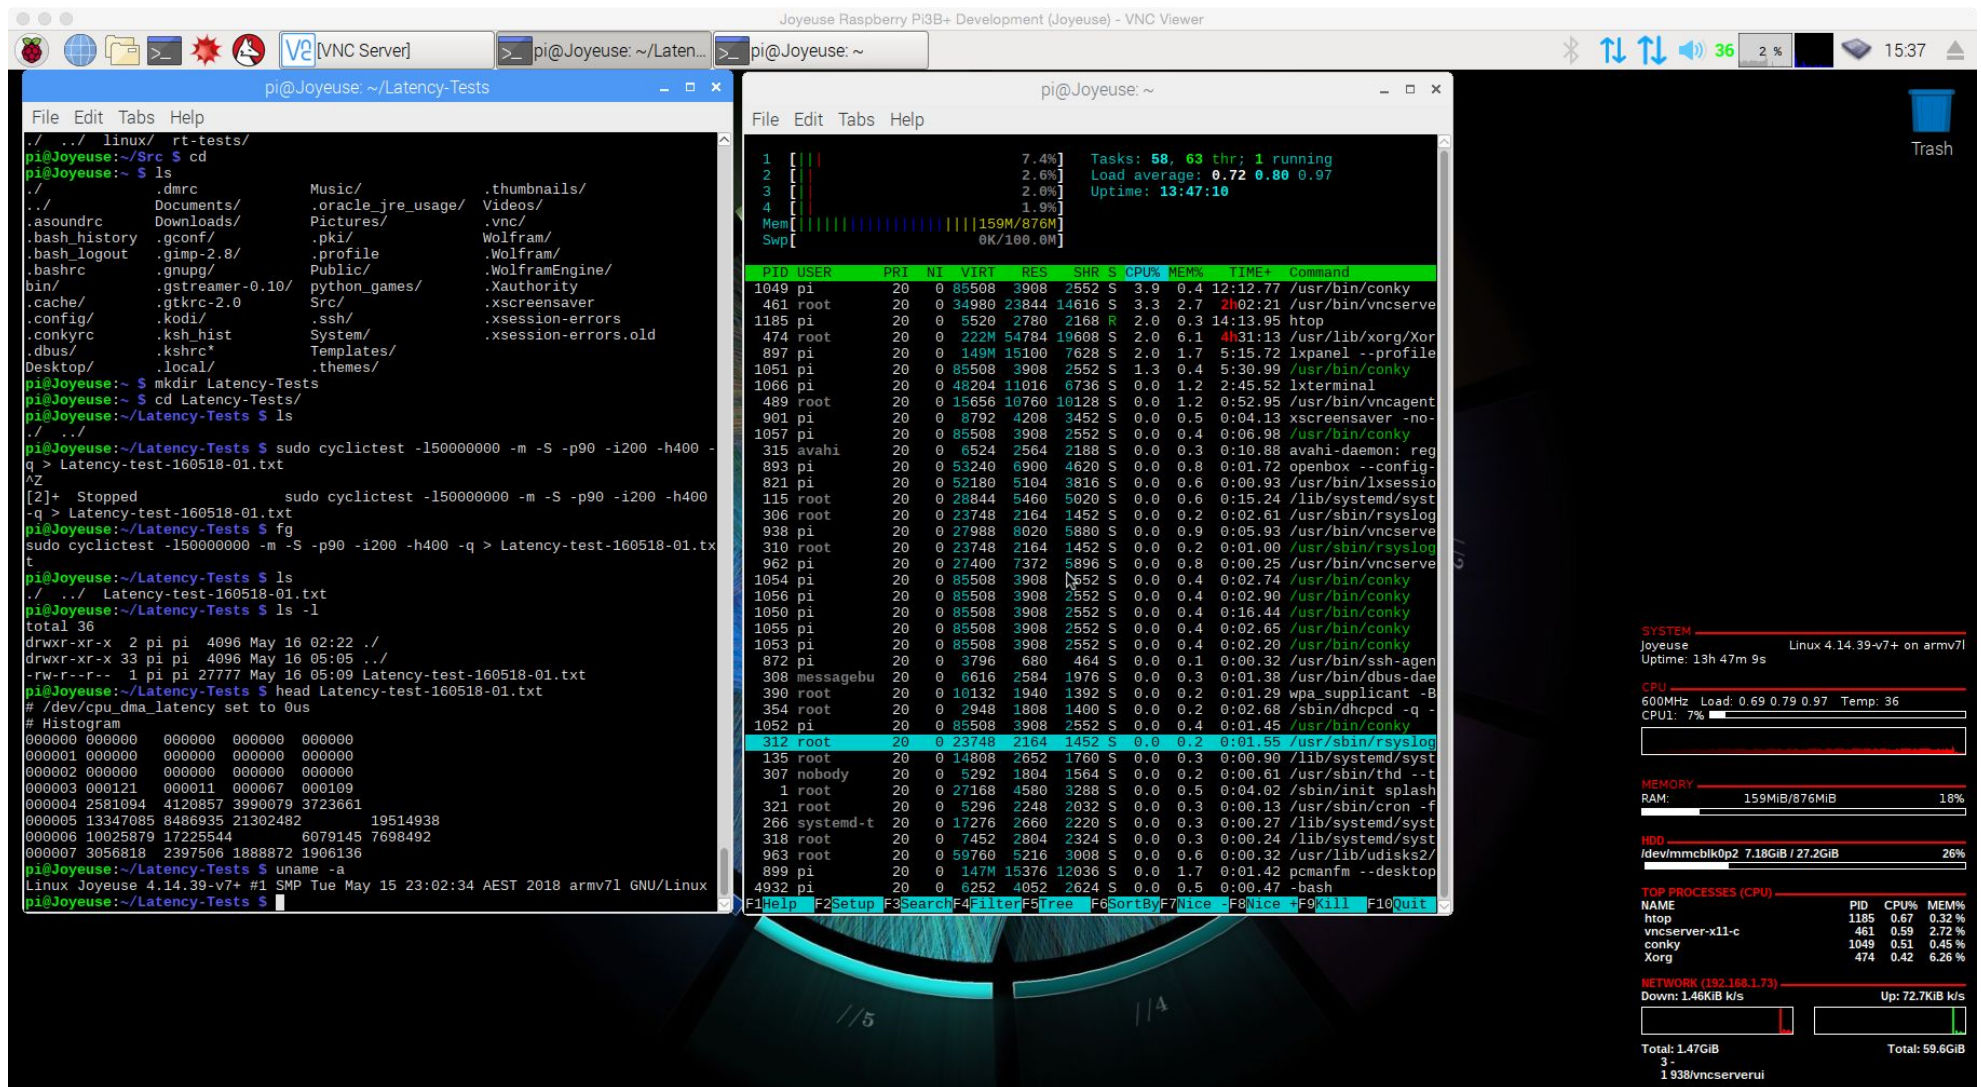


REAL-TIME SCHEDULING ALGORITHMS

PRIORITY DRIVEN SCHEDULING (1)



REAL-TIME LATENCY TESTING (RPI3B RTAI)



REAL-TIME LATENCY TESTING (RPI3B RTAI)

