



MONASH University

Information Technology

FIT2100 Operating Systems

THREADS AND CONCURRENCY

Week 7

Semester 2 2024

(Reading: Tanenbaum, Chapter 2)

Dr Charith Jayasekara

Faculty of Information Technology

© 2024 Monash University

OUTLINE

- The Thread Model
- Thread Usage
- Thread Implementations
- POSIX Threads
- Concurrency Problems
- Race Condition
- Mutual Exclusion

LEARNING OUTCOMES

- Upon the completion of Week 7, you should be able to:
 - Understand the **distinction between process and thread**
 - Describe the basic thread models
 - Explain the difference between **user-level threads** and **kernel-level threads**
 - Discuss thread management in Unix/Linux
 - Discuss the problems related to **concurrency**
 - Understand the concept of a **race condition**
 - Understand the concept of **critical regions**
 - Describe the **mutual exclusion** requirements

THE THREAD MODEL



THREADS

THE CONCEPT OF PROCESSES (REVISIT)

- The unit of **resource allocation** and a unit of **protection**.
- A (virtual) address space that holds the process image.
- Protected access to:
 - **Processor(s)**
 - **Other processes**
 - **Files**
 - **I/O resources**

THREADS

WHAT IS A THREAD?

- The unit of dispatching for execution is referred to as:
 - Thread or
 - Lightweight process
- The unit of resource ownership is referred to as:
 - Process or task
- **Multi-threading:**
 - The ability of an OS to support multiple concurrent paths of execution within a single process
 - **1 process : multiple threads of execution.**

THREADS


THREADS WITHIN A PROCESS

- ❑ Different part of a program may do different things and they can be executed **concurrently** to **improve response time** (or completion time).
 - Example: one thread may do a processor-bound task like rendering an image, while another thread responds to user interaction in the same program.
- ❑ If there is an interaction between different parts of the programs — **concurrency control** need to be applied.
- ❑ Example: accessing and modifying a common variable — **mutual exclusion** need to be satisfied.

THE THREAD MODEL

SINGLE-THREADED APPROACH


- ❑ A single thread of execution per process.
- ❑ The concept of a thread is not recognised — referred as a **single-threaded** approach.
- ❑ Example: MS-DOS, Windows 3.1



THE THREAD MODEL

MULTI-THREADED APPROACH

- ❑ One process with multiple threads.
- ❑ Example: Windows, Linux.




THE THREAD MODEL

PER-PROCESS vs PER-THREAD ITEMS

Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

- The first column lists some items *shared* by all threads in a process.
- The second one lists some items *private* to each thread.




Each thread has its own stack.

THREAD USAGE

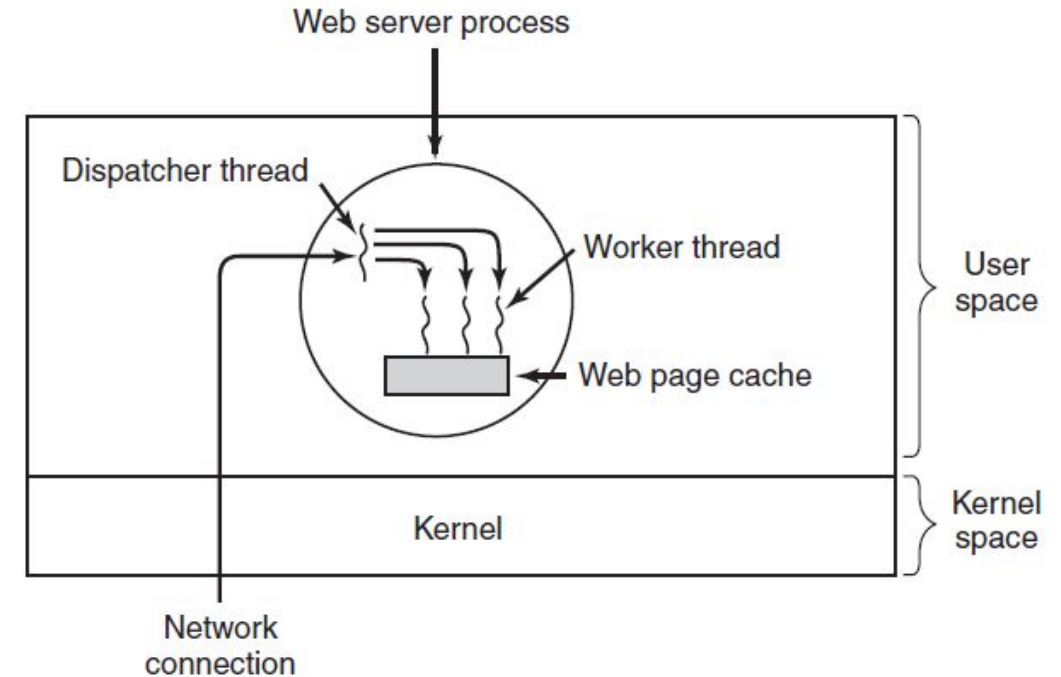


THREAD USAGE

EXAMPLES OF MULTI-THREADED APPLICATIONS



A word processor with three threads.



A multithreaded Web server.

THREAD USAGE

MULTI-THREADED WEB SERVER IMPLEMENTATION

```
while (TRUE) {  
    get_next_request(&buf);  
    handoff_work(&buf);  
}
```

(a)

```
while (TRUE) {  
    wait_for_work(&buf)  
    look_for_page_in_cache(&buf, &page);  
    if (page_not_in_cache(&page))  
        read_page_from_disk(&buf, &page);  
    return_page(&page);  
}
```

(b)

A rough outline of the code for a multi-threaded web server
(a) Dispatcher thread. (b) Worker thread.

THREAD USAGE

BENEFITS OF USING THREADS

- ❑ **Speed:** Takes less time to create a new thread than a process.
- ❑ **Speed:** Takes less time to terminate a thread than a process.
- ❑ **Speed:** Switching between two threads takes less time than switching between processes.
- ❑ **Speed:** Communication between threads within the same process is faster than communication between different processes.


DRAWBACKS OF USING THREADS

- ❑ **Reliability:** A bug in one thread can disable multiple threads, or all threads.
- ❑ **Security:** If one thread is compromised, usually all threads are compromised.
- ❑ **Consistency:** Concurrency problems may arise with shared data structures
- ❑ **Portability:** Support for multithreading depends on CPU type, and performance gains may not be portable across different CPUs even of the same nominal architecture


THREAD USAGE

BENEFITS OF USING THREADS: REMOTE PROCEDURE CALL (RPC) EXAMPLE

- ❑ For a single-threaded approach, 2 RPC requests are made one after another by Process 1 (Fig. a).
- ❑ Process 1 is completely blocked while waiting for a response to arrive from the server (Fig. a).
- ❑ For a multi-threaded approach, Thread A sends the first RPC request and gets blocked (Fig. b).
- ❑ Thread B gets some CPU time and sends the second RPC request (Fig. b).
- ❑ The process completes quickly as the second RPC response arrives shortly after the first RPC response. (Fig. b).
- ❑ **Assumption:** Server is also multi-threaded. OR Different servers handled the RPC requests.



(a) RPC using single thread



(b) RPC using one thread per server (on a uniprocessor)

- ▨ Blocked, waiting for response to RPC
- Blocked, waiting for processor, which is in use by Thread B
- Running

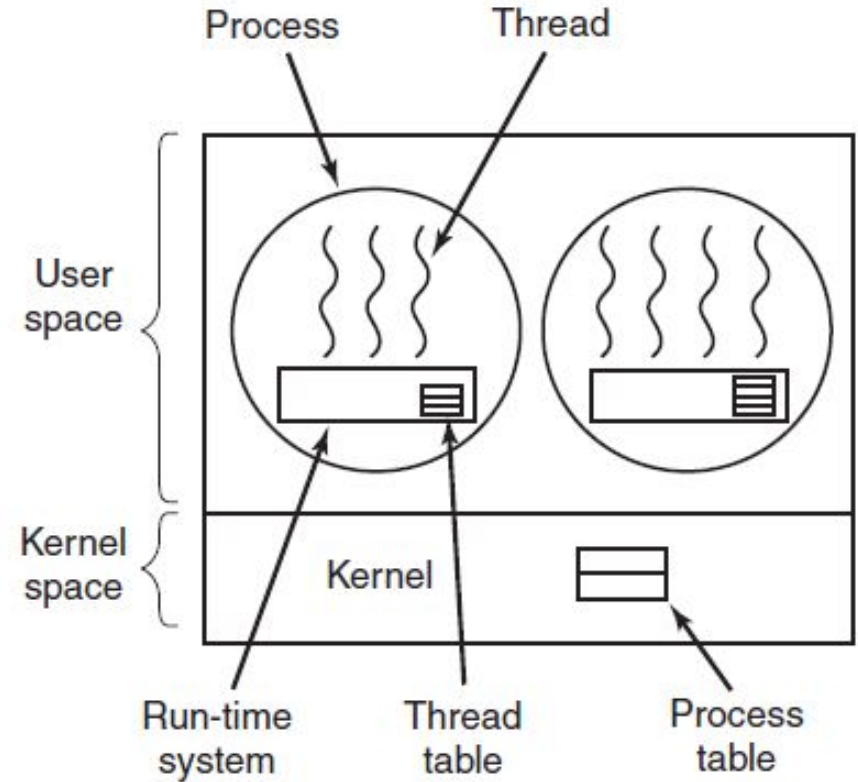
THREAD IMPLEMENTATIONS



THREAD IMPLEMENTATIONS

USER LEVEL THREADS (ULTs)

- ❑ All thread management is done by the application.
- ❑ The kernel is not aware of the existence of threads.
- ❑ Any application can be programmed to be multi-threaded by using a **threads library**.
 - Even if OS does not support threads.



THREAD IMPLEMENTATIONS

USER LEVEL THREADS (ULTs)

- **Advantages:**

- ❑ Thread switching does not require kernel mode privileges.
- ❑ Scheduling can be application specific.
- ❑ ULTs can run on any Operating Systems.


- **Disadvantages:**

- ❑ In a typical OS, many system calls are *blocking*.
- ❑ When a ULT executes a system call, not only is that thread gets blocked, but all of the threads within the process are also blocked.
- ❑ In a pure ULT strategy, a multi-threaded application cannot take the full advantage of multiprocessing.

THREAD IMPLEMENTATIONS

KERNEL LEVEL THREADS (KLTs)

- ❑ Thread management is done by the kernel.
- ❑ No thread management is done by the application — through API to the kernel thread facility.
- ❑ Example: Windows, Linux



THREAD IMPLEMENTATIONS

KERNEL LEVEL THREADS (KLTs)

- **Advantages:**


- ❑ The kernel can simultaneously schedule multiple threads from the same process on **multiple processors**.
- ❑ If one thread in a process is blocked, the kernel can schedule another thread of the same process.
- ❑ The kernel routines can also be multi-threaded.

- **Disadvantages:**

- ❑ The transfer of control from one thread to another thread within the same process requires a **mode switch** to the kernel.
 - Some overhead here.

THREAD IMPLEMENTATIONS


HYBRID IMPLEMENTATIONS



Multiplexing user-level threads
onto kernel-level threads.

THREAD IMPLEMENTATIONS

POP UP THREADS



Creation of a new thread when a message arrives:

(a) Before the message arrives. (b) After the message arrives.

POSIX THREADS (PTHREADS)



POSIX THREADS

PTHREADS LIBRARY

- ❑ Threads in Linux are known as **pthread**s. Managed through a separate API.
- ❑ The **pthread** library must be included and linked into the program in order to use threads.
 - `#include <pthread.h>`
 - Add `-lpthread` to the end of the **gcc** command to **link** the program against the pthread library
- ❑ `pthread_create()` – spawn a new thread
- ❑ `pthread_join()` – wait for another thread to terminate

POSIX THREADS

PTHREADS EXAMPLE PROGRAM

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

#define NUMBER_OF_THREADS 10

void *print_hello_world(void *tid)
{
    /* This function prints the thread's identifier and then exits. */
    printf("Hello World. Greetings from thread %d\n", tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[])
{
    /* The main program creates 10 threads and then exits. */
    pthread_t threads[NUMBER_OF_THREADS];
    int status, i;

    for(i=0; i < NUMBER_OF_THREADS; i++) {
        printf("Main here. Creating thread %d\n", i);
        status = pthread_create(&threads[i], NULL, print_hello_world, (void *)i);
```



POSIX THREADS

PTHREADS EXAMPLE PROGRAM

```
int status, i;  
  
for(i=0; i < NUMBER_OF_THREADS; i++) {  
    printf("Main here. Creating thread %d\n", i);  
    status = pthread_create(&threads[i], NULL, print_hello_world, (void *)i);  
  
    if (status != 0) {  
        printf("Oops. pthread_create returned error code %d\n", status);  
        exit(-1);  
    }  
}  
exit(NULL);  
}
```

CONCURRENCY PROBLEMS



CONCURRENCY PROBLEMS

WHY CONCURRENCY PROBLEMS OCCUR

- ❑ Processes that are working together may share some common storage or memory that each one can read and write to communicate between themselves.
- ❑ Processes need to access a shared resource in the system.
- ❑ Concurrent access to shared data may result in data inconsistency.
- ❑ Maintaining **data consistency** requires mechanisms to ensure the orderly execution of cooperating processes.
- ❑ Difficult to locate programming errors — results are non-deterministic and not reproducible.
- ❑ The problem exists in both multiprogramming on uni-processor and multi-processors.

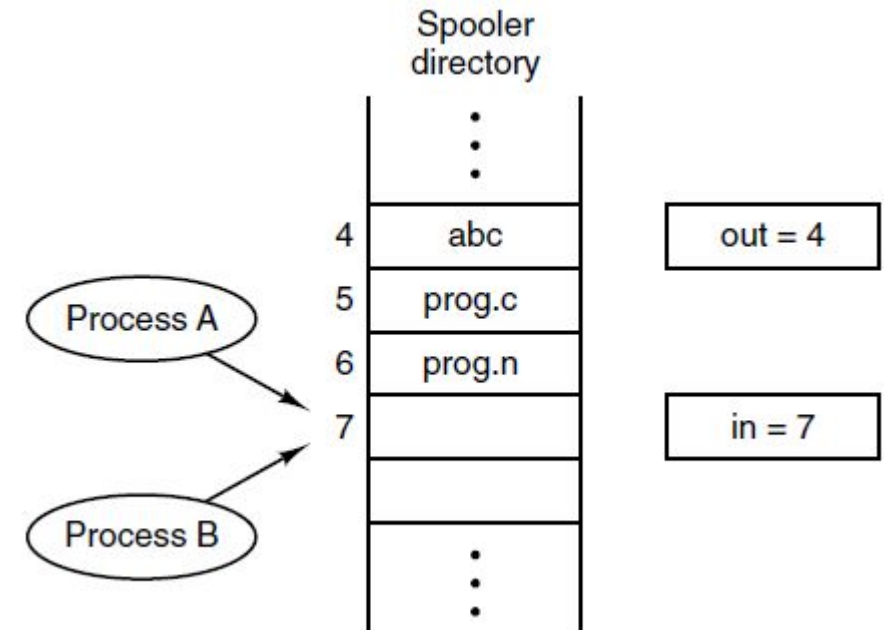
RACE CONDITIONS



RACE CONDITIONS

WHAT IS RACE CONDITION

- ❑ **A Race Condition** is a situation where two or more processes are reading or writing some shared data and the final result depends on who runs when.
- ❑ Occurs when multiple processes or threads read and write data items.
- ❑ Final result depends on the **order of execution**
 - **"Loser"** of the race is the process that updates last and will determine the final value of the variable
- ❑ To prevent race conditions, concurrent processes or threads must be **synchronised**.



Two processes want to access shared memory at the same time.

RACE CONDITION

A RACE CONDITION IN UNIPROCESSOR MULTIPROGRAMMING

- Consider the following procedure:

```
void echo()  
{  
    char_in = getchar();  
    char_out = char_in;  
    putchar(char_out);  
}
```

- Read a character from the keyboard and store in `char_in`.
- Transfer to `char_out` before being sent for display.
- Consider two different applications — `Process A` and `Process B` — make a call to this procedure.

RACE CONDITION

RACE CONDITION IN UNIPROCESSOR MULTIPROGRAMMING

- Consider the following procedure:

```
void echo()  
{  
    char_in = getchar();  
    char_out = char_in;  
    putchar(char_out);  
}
```

- Process A invokes `echo()` and is interrupted immediately after `getchar` returns its value and stores it in `char_in` (e.g. `x`).
- Process B is activated and invokes `echo()` and read a char (e.g. `y`) and runs to completion of the procedure.
- When process A resumes the value of `x` has been overwritten in `char_in` by process B and therefore its value `x` is lost.

RACE CONDITION

RACE CONDITION IN UNIPROCESSOR MULTIPROGRAMMING

- Process A:

```
char_in = getchar();
```

```
char_out = char_in;  
putchar(char_out);
```

- Process B:

```
char_in = getchar();  
char_out = char_in;  
putchar(char_out);
```



TIME

MUTUAL EXCLUSION



MUTUAL EXCLUSION

WHAT IS MUTUAL EXCLUSION (MUTEX)

- Suppose n processes all **competing** to use some shared data.
- Each process has a code segment — **critical region** — where the shared data is accessed or manipulated.
- **Mutual exclusion** refers to the requirement of ensuring that no two concurrent processes are in their critical region at the same time.
- Mutual Exclusion is a basic requirement in concurrency control, to prevent **race conditions**.
- Ensure that when one process is executing in its critical region, **no other process is allowed in its critical region**.

MUTUAL EXCLUSION

REQUIREMENTS TO AVOID RACE CONDITIONS

- ❑ No two processes may be simultaneously inside their critical regions.
- ❑ No assumptions may be made about speeds or the number of CPUs.
- ❑ No process running outside its critical region may block other processes.
- ❑ No process should have to wait forever to enter its critical region.

MUTUAL EXCLUSION

PREVENTING A RACE CONDITION BY USING MUTUAL EXCLUSION

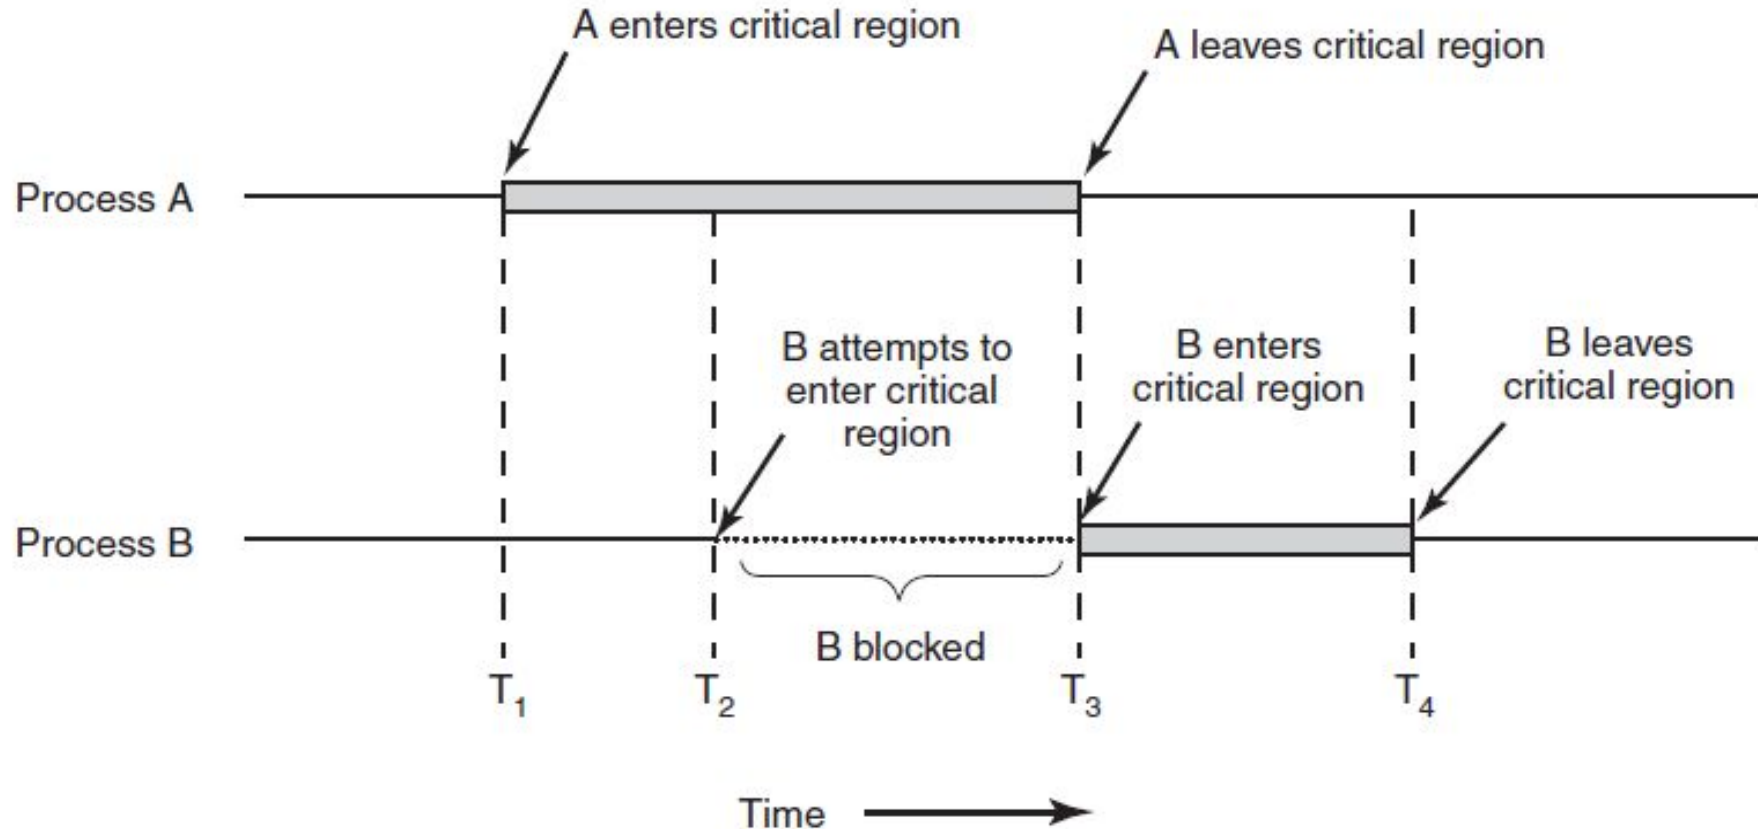
- Consider the following procedure:

```
void echo()  
{  
    char_in = getchar();  
    char_out = char_in;  
    putchar(char_out);  
}
```

- Assume **only one process at a time** to invoke and be in the **echo** procedure (**the whole echo procedure is the critical region**)
- Process A invokes the **echo** procedure and is interrupted immediately after **getchar** returns its value and stores it in **char_in** (e.g. x).
- Process B is activated and invokes **echo** procedure and since the **echo** procedure is used by process A, **process B is blocked from further execution**.
- At some later time, process A is resumed and completes the execution of **echo** and the proper input character will be displayed.
- When process A exits **echo**, **this removes the block on process B**.
- When process B is later resumed, the echo procedure is successfully invoked.

MUTUAL EXCLUSION

PREVENTING A RACE CONDITION BY USING MUTUAL EXCLUSION



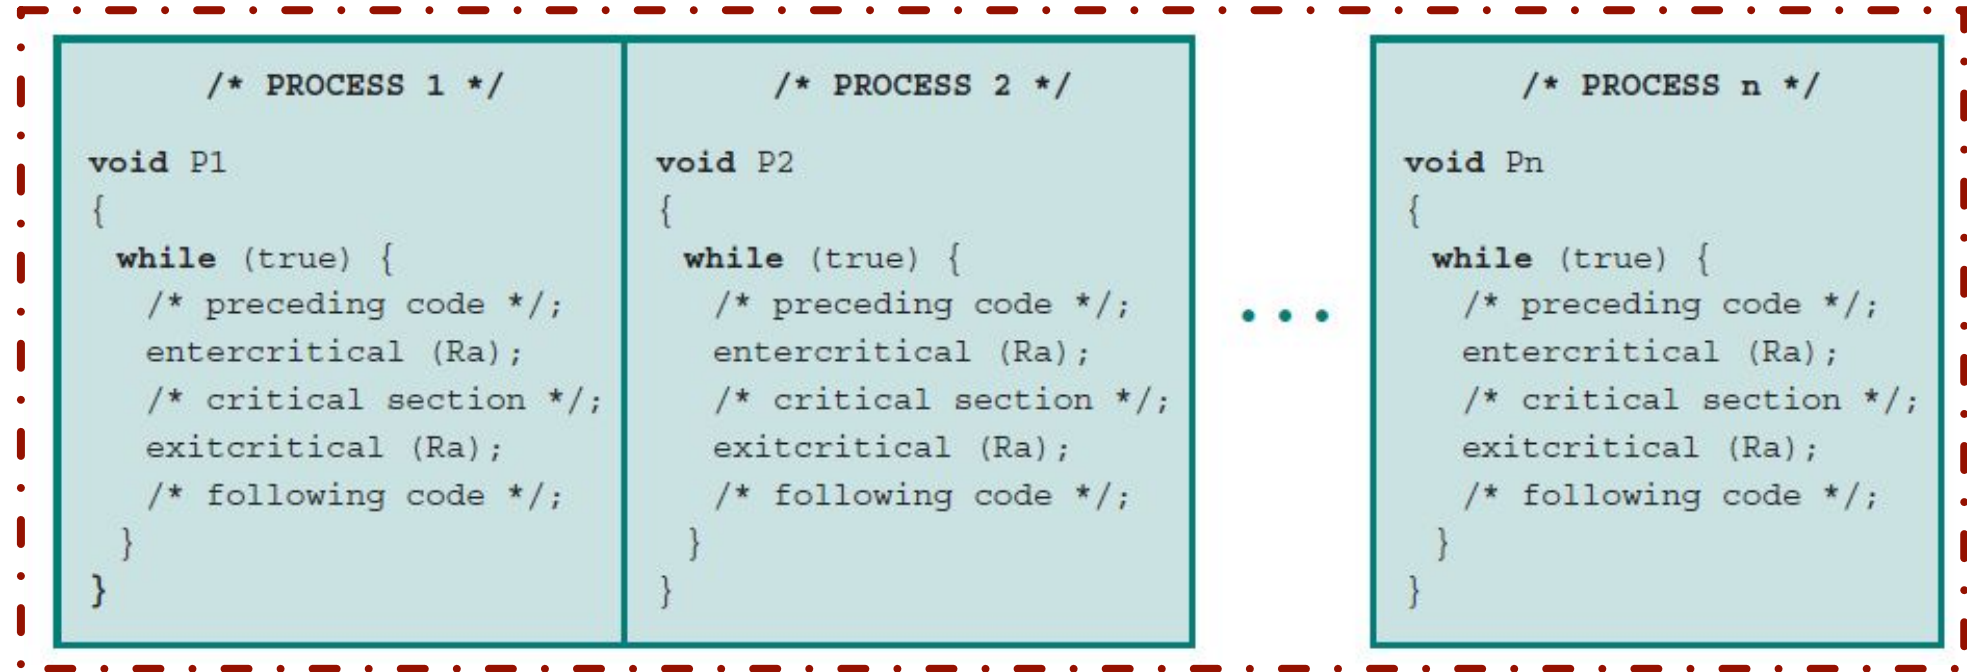
MUTUAL EXCLUSION

RACE CONDITION IN MULTIPROCESSOR MULTIPROGRAMMING

- ❑ Same problem arises even when the processes — A and B runs on **different processors** accessing unprotected shared variables.
- ❑ The solution outlined in the previous slides can work here.
- ❑ **Protecting and controlling access to shared resources are critical.**

MUTUAL EXCLUSION

EXAMPLE



To enforce mutual exclusion, two function are provided: *entercritical* and *exitcritical* with the resource as the argument.

MUTUAL EXCLUSION

RACE CONDITIONS WITH MULTIPLE SHARED DATA RESOURCES

- ☐ The same problem exists even when processes access more than one shared resource.
- ☐ Processes must **cooperate** to ensure the shared data are properly managed.
- ☐ Control mechanisms are needed to ensure the **integrity** of the shared data.

MUTUAL EXCLUSION

RACE CONDITION EXAMPLE WITH MULTIPLE SHARED DATA RESOURCES

▪ Process A

```
a = a + 1 ;  
b = b + 1 ;
```

▪ Process B

```
b = 2 * b ;  
a = 2 * a ;
```

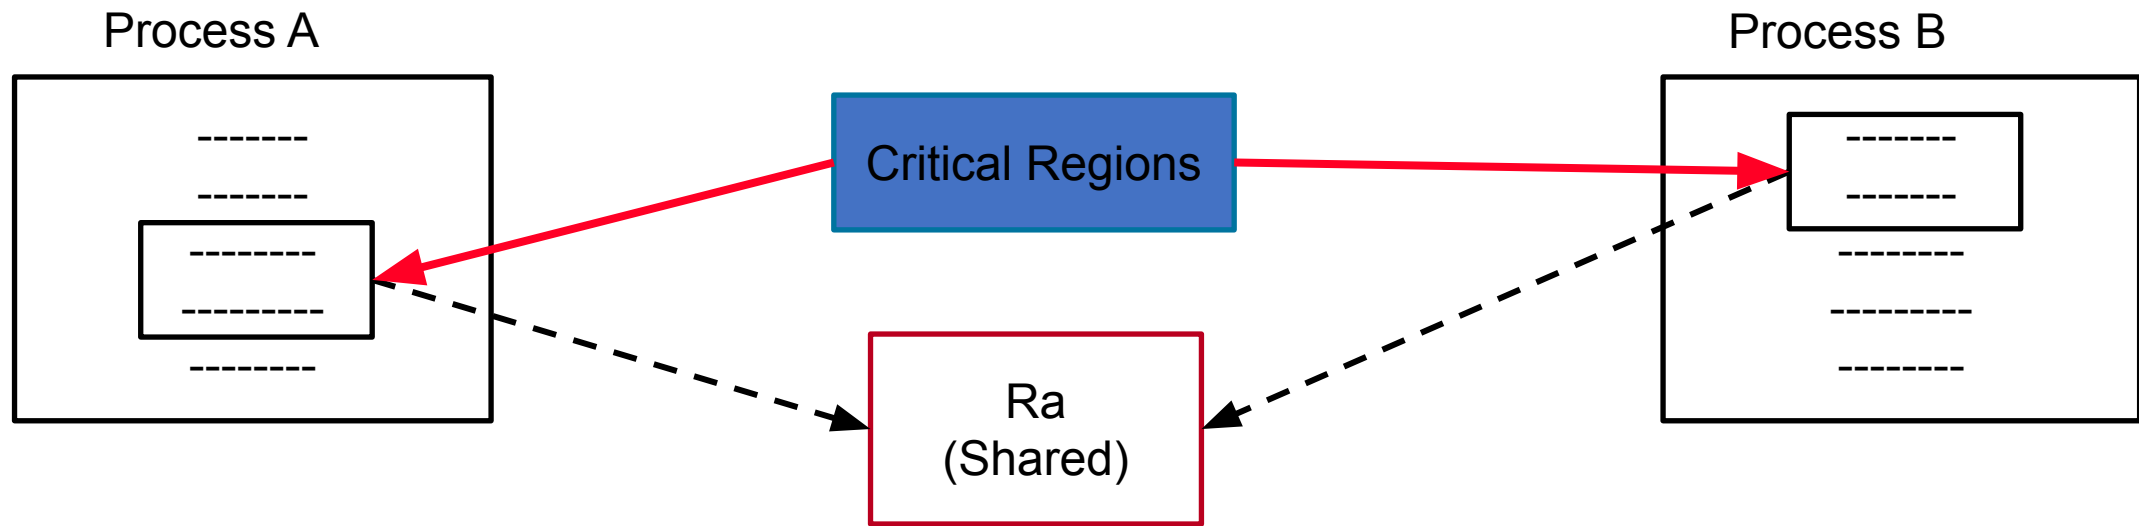
- Assuming that **a = b** at the beginning, and consider the following concurrent execution sequence:

```
a = a + 1 ;    /* {PA} */  
b = 2 * b ;    /* {PB} */  
b = b + 1 ;    /* {PA} */  
a = 2 * a ;    /* {PB} */
```

- At the end of this execution, the condition **a = b** *no longer holds!*
- **Solution:** access to both a and b should be enclosed in the same critical region.

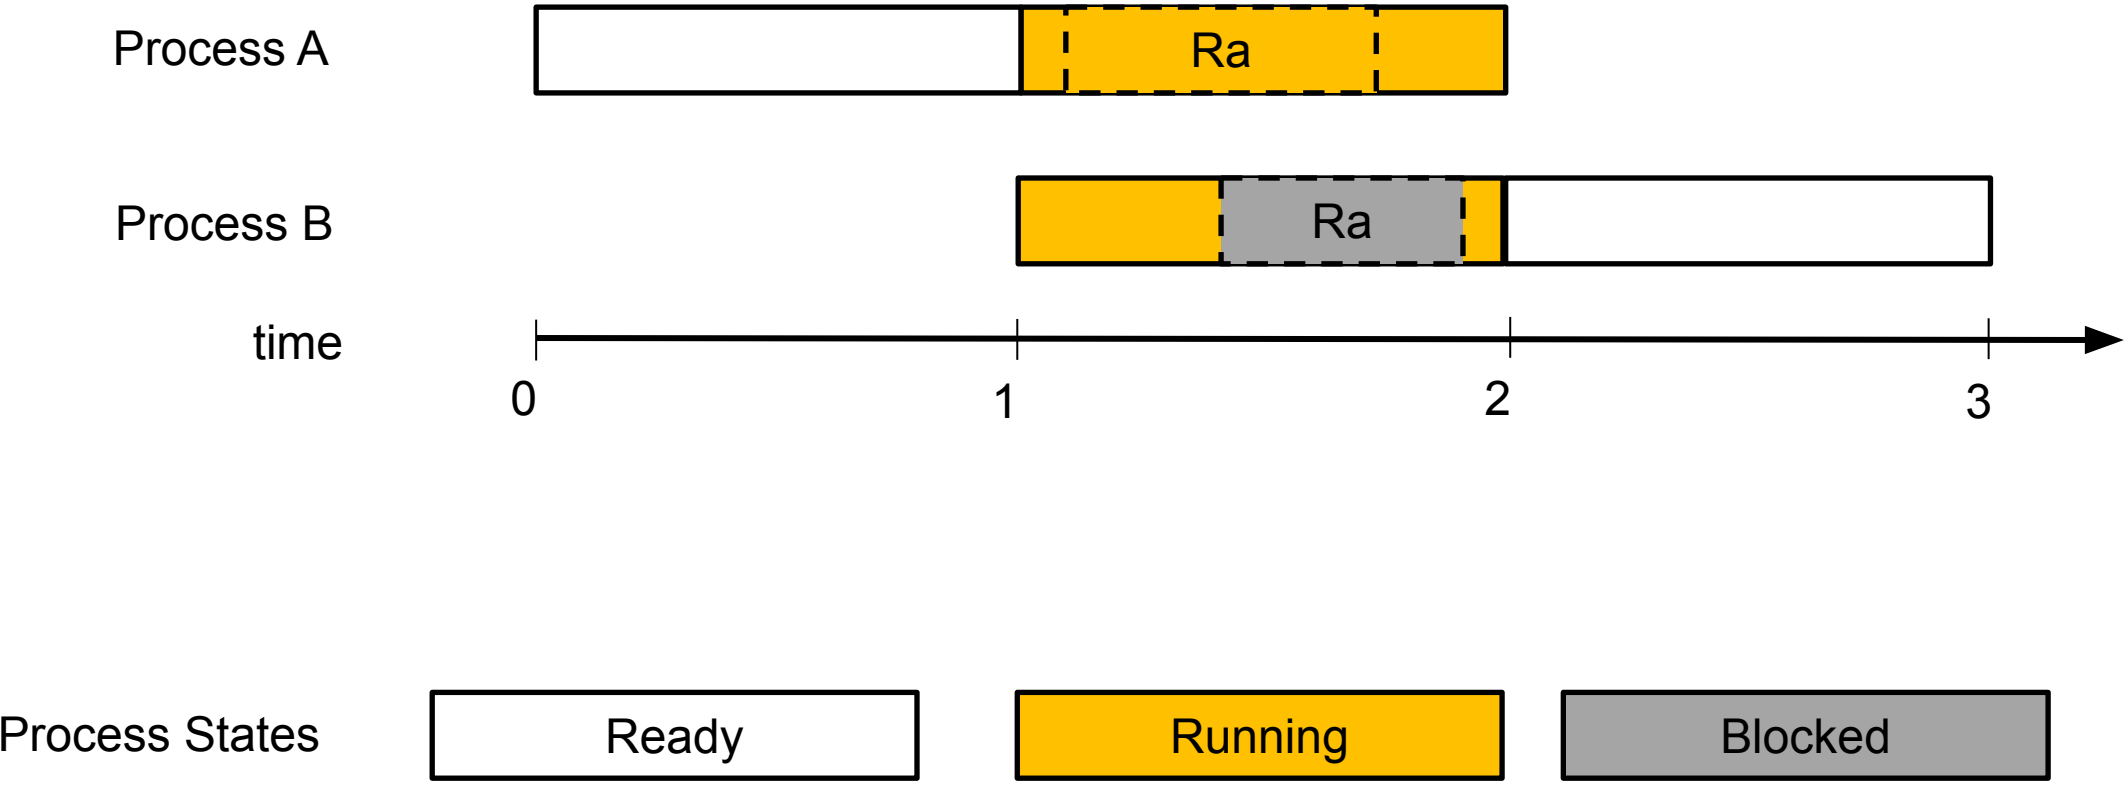
RACE CONDITION

REVISIT





RACE CONDITION

REVISIT



RACE CONDITION

RACE CONDITION PREVENTION WITH MUTUAL EXCLUSION



CONCURRENCY PROBLEMS

MORE CONCURRENCY PROBLEMS

- ❑ **Mutual Exclusion causes additional concurrency problems!!!**
- ❑ **Deadlock**: two or more processes are waiting indefinitely for the other processes to release the system resources.
- ❑ **Starvation**: indefinite blocking of a process.

Summary

- ❑ So far we have discussed
 - Introduction to threads
 - Thread Model
 - Thread implementations
 - Concurrency Problems
- ❑ Next week
 - Advanced Concurrency
- ❑ Reading
 - Tanenbaum, Chapter 2 (4th Edition)
 - Stallings, Chapter 7 (7th Edition)

